



Technische
Universität
Braunschweig



Dynamic parallelism using CUDA

The shallow water equation on an adaptive grid


Marc Kassubeck, Torsten Thoben, 18. März 2015

Inhalt

- **The task**
- **Dynamic parallelism using CUDA**
 - Implementation of DP with CUDA
 - DP and Synchronisation
 - Recursion Depth and Device Limits
- **The Problem**
 - First try
 - Back to Arrays
 - Results and discussion

The Task

- Extension of an existing shallow water solver
- With dynamic parallelism (DP)



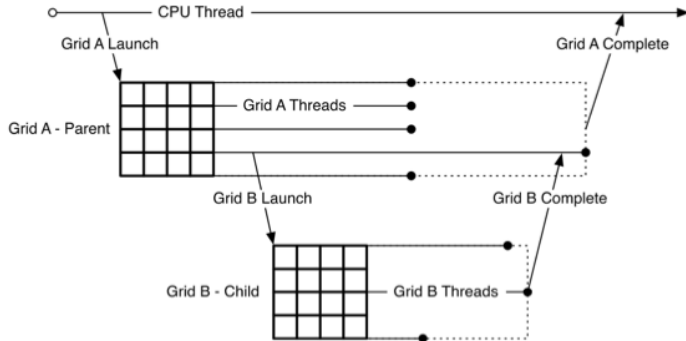
Implementation of DP with CUDA

- DP in CUDA means calling a kernel inside a kernel
- So every thread calls a new grid

```
main(){  
    ...  
    func <<< g,b >>> (depth);  
}  
  
__global__ void func(int depth){  
    ...  
    if(depth<2)  
        func <<< g, b >>> (depth+1);  
}
```

DP and Synchronisation

- Father-kernel finished if his subgrid is finished but can continue working after the subkernel call
- To synchronize between one thread and his subgrids use `cudaDeviceSynchronize()` (\neq `__syncthreads()`)

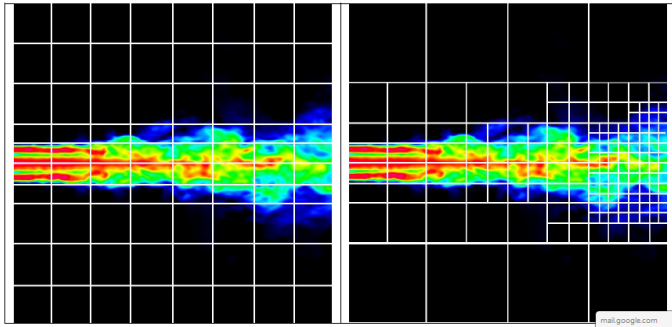


Recursion Depth and Device Limits

- Nesting depth (hardware-limit is 24)
 - Need to reserve a buffer for running/suspended or launching kernels
 - `cudaDeviceSetLimit(cudaLimitDevRuntimePendingLaunchCount, x)`
 - Default is set to 2048
- Synchronization depth
 - `cudaDeviceLimit(cudaLimitDevRuntimeSyncDepth, x)`
 - Default is set to 2

Problem: Implementation of the datastructure

- Datastructure: Forest (many Trees)
- Limit to coarseness is the `forestsize` (number of trees)
- Limit to getting fine end up to the maximum recursions
- Trees: eg Quadtrees, Octrees...




GPU Datastructure using pointer

```
class LeafElem : public TreeElem
{
public:
    float value;
    __device__ __host__ LeafElem();
    virtual __device__ __host__ ~LeafElem(){};
    virtual __device__ __host__ bool isLeaf()
    {
        return true;
    }
};
```


GPU Datastructure using pointer

```
class BranchElem : public TreeElem
{
public:
    int nx;
    int ny;
    int depth;
    TreeElem** children;
    __device__ __host__ BranchElem(int nx, int ny, int depth)
    virtual __device__ __host__ ~BranchElem();
    virtual __device__ __host__ bool isLeaf()
    {
        return false;
    }
};
```

But this fails




- Kernel calls take too long to execute (> 7 seconds)
- Possible causes:
 - Exponentially growing trees
 - Overhead to launch subgrids
 - No coalesced memory access with pointers of this data structure

Storing the tree-structure in an Integer array


0	0	0	0	1	1	2	2
0	0	0	0	1	1	2	2
0	0	0	0	1	1	1	1
0	0	0	0	1	1	1	1
0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0

- Integer values denote the level in the tree
- The higher the value, the finer the grid will be in this area
- Solution values are averaged and stored in all corresponding 'supercells'
- Allows for better coalesced memory access

Profiling



Results



Problems

- Unfortunately this is still the solution using the original solver
- Some problems surfaced, when implementing the dynamic version
- Implementing the memory using CUDA Unified memory:
 - On the one hand no `cudaMemcpy`
 - On the other hand many calls to `cudaDeviceSynchronize`
 - Too much data transfer CPU \leftrightarrow GPU
 - Too slow using our current implementation
- Problems in the solver code
 - Solutions may contain ∞ or `#INDEFINITE`
 - Maybe stability issue
 - Maybe bug in solver code